

ITCS 4152/5152: Speed Limit Sign Detection

Group 11: Thomas Carr, Kyle Ward, Payne
Miller, Jay Yadav, Makalia Vang




Introductions



- Payne Miller
 - Preparation of Data & Research
 - Assisted with Backend Demo
 - Presentation Creation



- Kyle Ward
 - Model Creation & Research
 - Presentation Creation



- Jay Yadav
 - Preparation of Data & Research
 - Assisted with Backend & Frontend Demo
 - Presentation Creation



Introductions



- Thomas Carr
 - Preparation of Data & Research
 - Assisted with Backend & Frontend Demo
 - Model Creation & Research



- Makaila Vang
 - Presentation Creation
 - Preparation of Data & Research



Problem and Motivation

- Speed Limit Detection
 - a. Sometimes the speed limit can switch suddenly and traditional GPS systems don't always have the right information stored (outdated.)
 - b. Speeding Tickets are expensive.
- There are two main customers for our project:
 - a. Autonomous vehicles
 - Need to know what speed it is allowed to do without relying on external data.
 - b. Map/GPS companies.
 - Can use Street View images in order to get the speed limit of roads.



Our Dataset

- We used 1,000 images of various quality, lighting, and speed limit values within our model.
- Obtained images via the Open Source platform [Kaggle](#).
- Labels follow the format: speedLimit{Value}_{id}.avi_image{imageNumber}.png
- Train/Test/Validation Split:
 - a. Train (70%): 767 images
 - b. Test (10%): 106 images
 - c. Validation(20%): 216 images



Our Dataset (cont.)

speedLimit25_1398812718.avi_image4.png



speedLimit25_1398991101.avi_image2.png



speedLimit25_1405384269.avi_image21.png



speedLimit25_1398812718.avi_image2.png



speedLimit25_1405361839.avi_image2.png



speedLimit30_1405371792.avi_image0.png



speedLimit25_1405384333.avi_image21.png



speedLimit30_1405359152.avi_image3.png



speedLimit25_1405044028.avi_image0.png



speedLimit30_1398813050 (1).avi_image5.png



speedLimit30_1405383641 (1).avi_image11.png



speedLimit25_1398812671.avi_image3.png



speedLimit40_1405118917.avi_image0.png



speedLimit35_1405362300.avi_image2.png



speedLimit40_1405372703.avi_image2.png



speedLimit35_1399493486.avi_image3.png



speedLimit40_1405118286.avi_image3.png



speedLimit30_1398991269.avi_image5.png



speedLimit40_1405106877.avi_image3.png



speedLimit35_1404947962.avi_image3.png



speedLimit25_1398812407.avi_image3.png



speedLimit25_1405384014.avi_image39.png



speedLimit30_1398993191.avi_image1.png



speedLimit30_1398991322.avi_image2.png



speedLimit40_1405111351.avi_image1.png



speedLimit30_1398991040.avi_image2.png



50-0.png



speedLimit30_1405383310.avi_image10.png



speedLimit25_1405466603.avi_image10.png



speedLimit30_1405383730.avi_image19.png





Our Methodology

- Our model is a CNN built using Tensorflow 2
 - a. Utilizes max pooling, average pooling, and upsampling layers for feature extraction
 - b. Adam optimizer with Categorical Cross-entropy loss function
 - c. Output is a probability distribution of possible classes (25,30,35,40,45,50mph)
- Outside-In approach:
 - Assumptions: Speed limit sign is square, numbers will always be contained within this square. Number will exist solely of lines and curves
 - First Conv-Pooling layer uses a max-pool pair for extracting the parts that define the lines of the rectangle and parts of the numbers
 - Second Conv-Pooling layer pair uses an average-pool to extract the contours and curvatures of the numbers
 - These are upsampled before feeding into the MLP part of the network
- We transformed the images into 64 x 64 pixels and applied a Grayscale filter.



Our Methodology - Second Model

- 3 Convolutions Before Any Pooling
 - In hopes of finding a new parameters missed by a simpler model
- Reshaped Images to 224x224x3
 - Used RGB instead of BW
- Added More Data Augmentation
 - Gaussian Blur
 - Rotation (30°)
 - X/Y Shifting (30px)
- Regularization
 - Batch Normalization
 - Dropout
 - 80/20 train/test split

Layer (type)	Output Shape	Param #
conv2d_44 (Conv2D)	(None, 224, 224, 16)	784
conv2d_45 (Conv2D)	(None, 224, 224, 16)	4112
conv2d_46 (Conv2D)	(None, 224, 224, 16)	4112
batch_normalization_18	(None, 224, 224, 16)	64
average_pooling2d_8	(None, 112, 112, 16)	0
conv2d_47 (Conv2D)	(None, 112, 112, 32)	4640
conv2d_48 (Conv2D)	(None, 112, 112, 32)	9248
conv2d_49 (Conv2D)	(None, 112, 112, 32)	9248
batch_normalization_19	(None, 112, 112, 32)	128
average_pooling2d_9	(None, 56, 56, 32)	0
conv2d_50 (Conv2D)	(None, 56, 56, 64)	18496
conv2d_51 (Conv2D)	(None, 56, 56, 64)	36928
conv2d_52 (Conv2D)	(None, 56, 56, 64)	36928
batch_normalization_20	(None, 56, 56, 64)	256
max_pooling2d_12	(None, 28, 28, 64)	0
conv2d_53 (Conv2D)	(None, 28, 28, 128)	73856
conv2d_54 (Conv2D)	(None, 28, 28, 128)	147584
conv2d_55 (Conv2D)	(None, 28, 28, 128)	147584
batch_normalization_21	(None, 28, 28, 128)	512
max_pooling2d_13	(None, 14, 14, 128)	0
flatten_5 (Flatten)	(None, 25088)	0
dense_25 (Dense)	(None, 256)	6422784
dropout_17 (Dropout)	(None, 256)	0
dense_26 (Dense)	(None, 128)	32896
dropout_18 (Dropout)	(None, 128)	0
dense_27 (Dense)	(None, 32)	4128
dropout_19 (Dropout)	(None, 32)	0
dense_28 (Dense)	(None, 16)	528
dense_29 (Dense)	(None, 7)	119

Total params: 6,954,935

Trainable params: 6,954,455

Non-trainable params: 480



Results

- Our best model achieved **90% accuracy**.
- From a development perspective:
 - a. Started with a base model of 85% accuracy
 - b. From there we iterated to 86% accuracy
 - c. Ultimately optimized up to 90% accuracy.
- We used the simple formula for Accuracy:
 - a. $\text{Accuracy} = \text{number of correct reads} / \text{total number of reads}$.
- Accuracy was categorical cross entropy.



Results (cont.)

Predicted:45
Actual:45



Predicted:35
Actual:30



Predicted:35
Actual:35



Predicted:35
Actual:35



Predicted:25
Actual:25



Predicted:45
Actual:45



Predicted:25
Actual:25



Predicted:25
Actual:25



Predicted:50
Actual:50



Predicted:35
Actual:35



Predicted:25
Actual:25



Predicted:25
Actual:25



Predicted:30
Actual:40



Predicted:30
Actual:30



Predicted:45
Actual:45



Predicted:25
Actual:25



Predicted:30
Actual:30



Predicted:50
Actual:50



Predicted:45
Actual:45



Predicted:25
Actual:30



Predicted:30
Actual:40



Predicted:25
Actual:25



Predicted:35
Actual:35



Predicted:30
Actual:30



Predicted:25
Actual:25



Predicted:25
Actual:25



Predicted:35
Actual:35



Predicted:25
Actual:25



Predicted:25
Actual:25



Predicted:25
Actual:25





Conclusions

- Our best model achieved **90% accuracy**.
- Faced some challenges
 - a. Scheduling between full-time students and full-time employees/part-time students.
 - b. We were pre-processing the images on our own, but then figured out a CNN would do it better.
 - c. Picking a topic of interest
 - d. Finding a large, diverse, dataset was a huge boon.



Live Demo

- Hosted our backend server on a personal server machine.
- Wanted to hit a stretch goal of a live system demonstration.
- speadle.web.app



Q & A

- Any questions?



Thank you!

- Thank you Professor Archit and TAs.